Cloud Computing Redaction Project

Michail Angelos Karvelas

Project Overview

The aim of this project is to establish a secure system that anonymizes sensitive documents, similar to a redaction service.

1. **The Goal**

Create a trusted system that protects privacy by anonymizing documents.
Automatically clear names, emails, and addresses from text.

2. **The Solution**

A simple service that receives text and sends back a fully cleaned PDF.
Powered by language models that find sensitive details with context.

3. **The Architecture**

Built for the cloud with containers and Kubernetes.
Ready for high demand with smart routing and traffic control.
Sensitive content is erased before any file is generated.

The Problem & Solution

- The Challenge: Handling sensitive documents (PII) securely is difficult. Manual redaction is slow and error-prone.
- The Solution: An automated Microservice Architecture that detects and permanently removes PII.
- Key Capabilities:
  - Intelligent Detection: Combines Regex (Email/Phone) and AI (Names/Locations).
  - Visual Redaction: Physically draws black boxes on PDFs (not just covering text).

Scalable: Runs on a self-healing Cloud Cluster (AKS).

System Architecture

- Frontend (Consumer):
User Interface for uploading files.

- Backend (Producer):
Python FastAPI service for OCR and NLP processing.

- Service Discovery:
Consul allows the Frontend to find the Backend dynamically.

System Architecture

Backend Explanation

Regex: Used for deterministic patterns (Emails, Phones). High speed, 100% accuracy for fixed formats.

NLP (spaCy): Used for context-aware entities (Names, Cities). Capable of distinguishing "Park" (Name) from "Park" (Location).

Why Hybrid?

Combining both covers the weaknesses of each, ensuring maximum accuracy.

System Architecture

Backend Code Screenshot

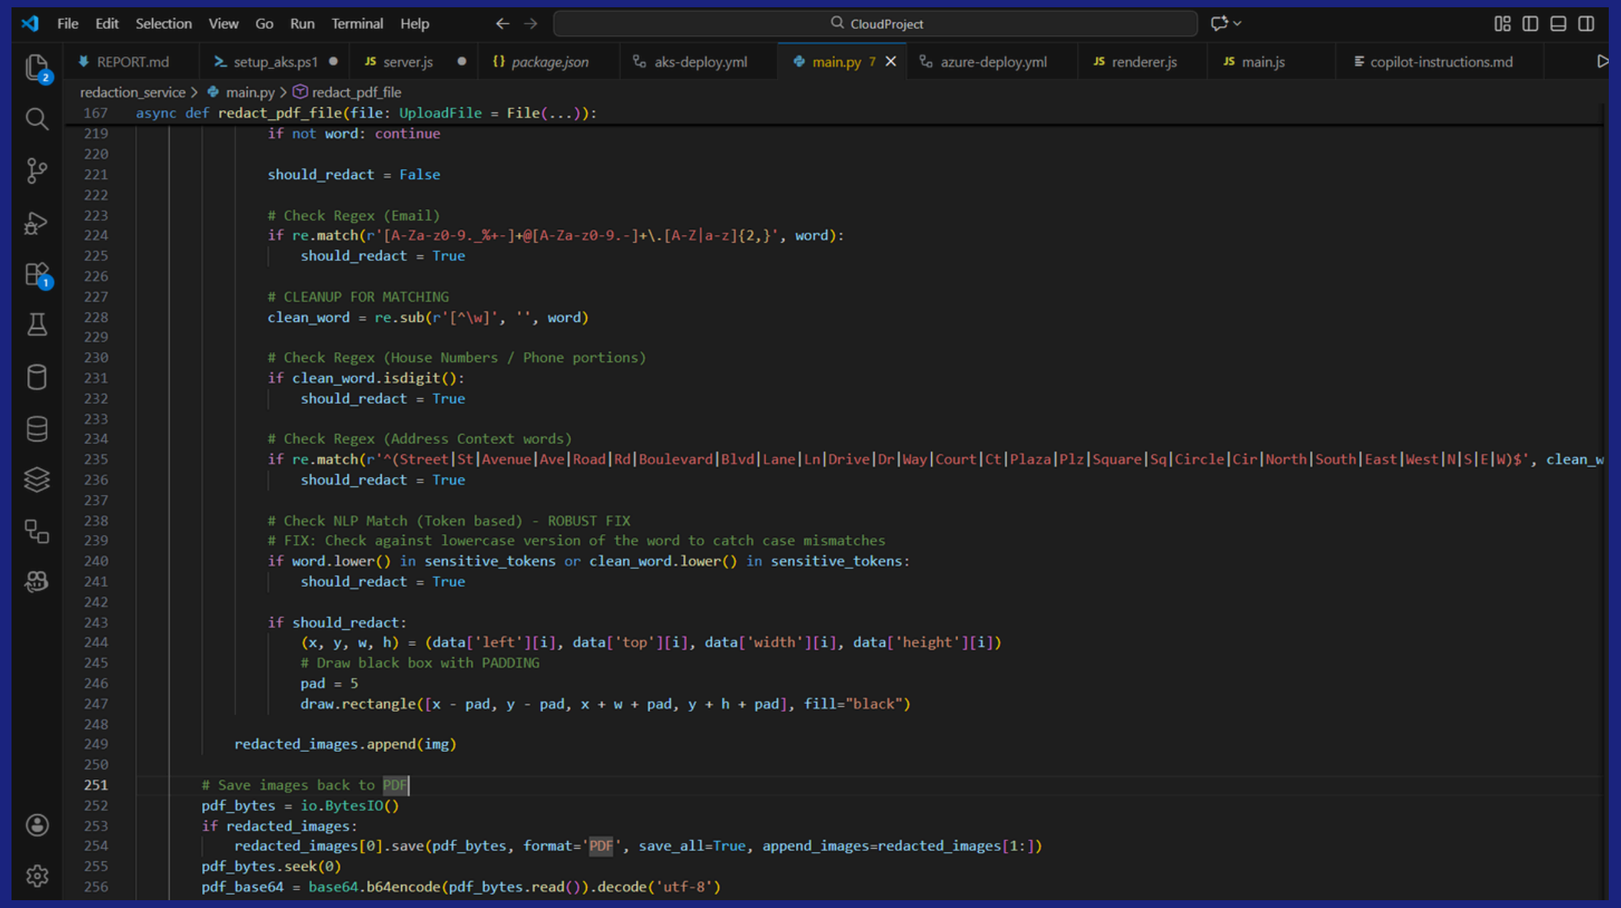

System Architecture

FrontEnd

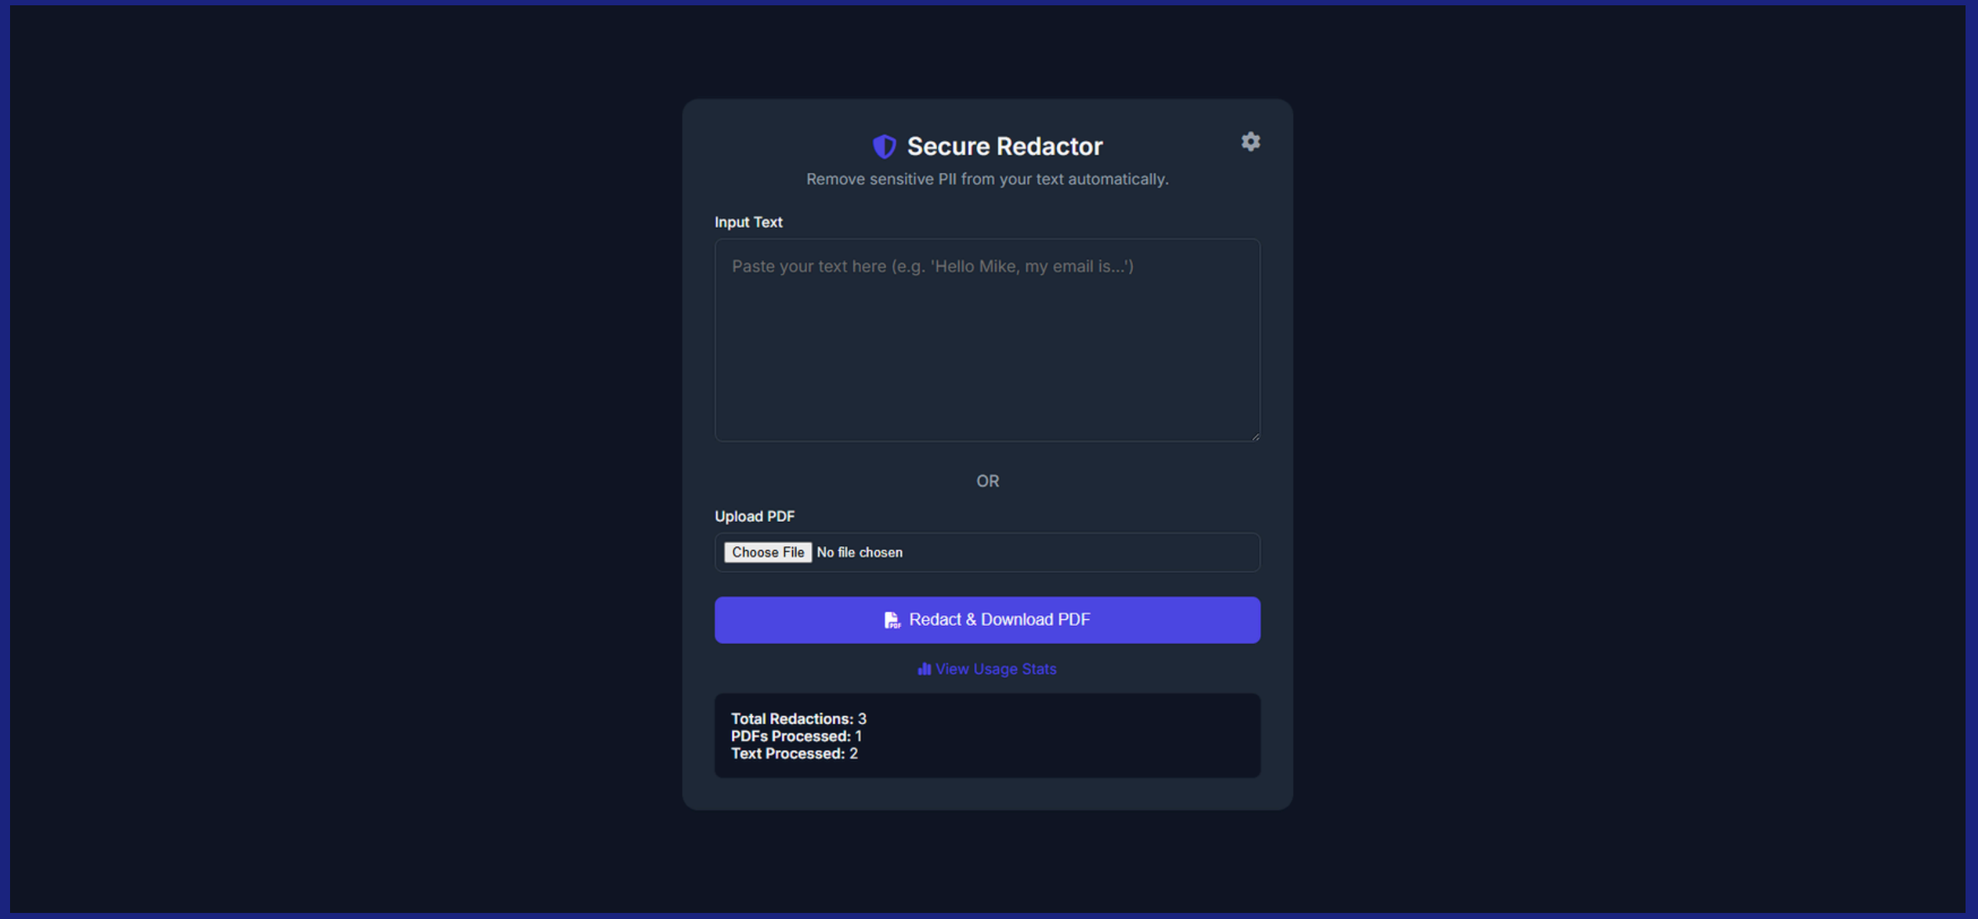

Kubernetes Infrastructure (AKS)

- Key Specs:
  - Cluster: Azure Kubernetes Service (AKS)
  - Nodes: 2 Nodes (Standard_B2s) for hardware redundancy
- Networking:
  - LoadBalancer: Exposes the Frontend to the public internet

ClusterIP: Keeps Backend communication private and secure

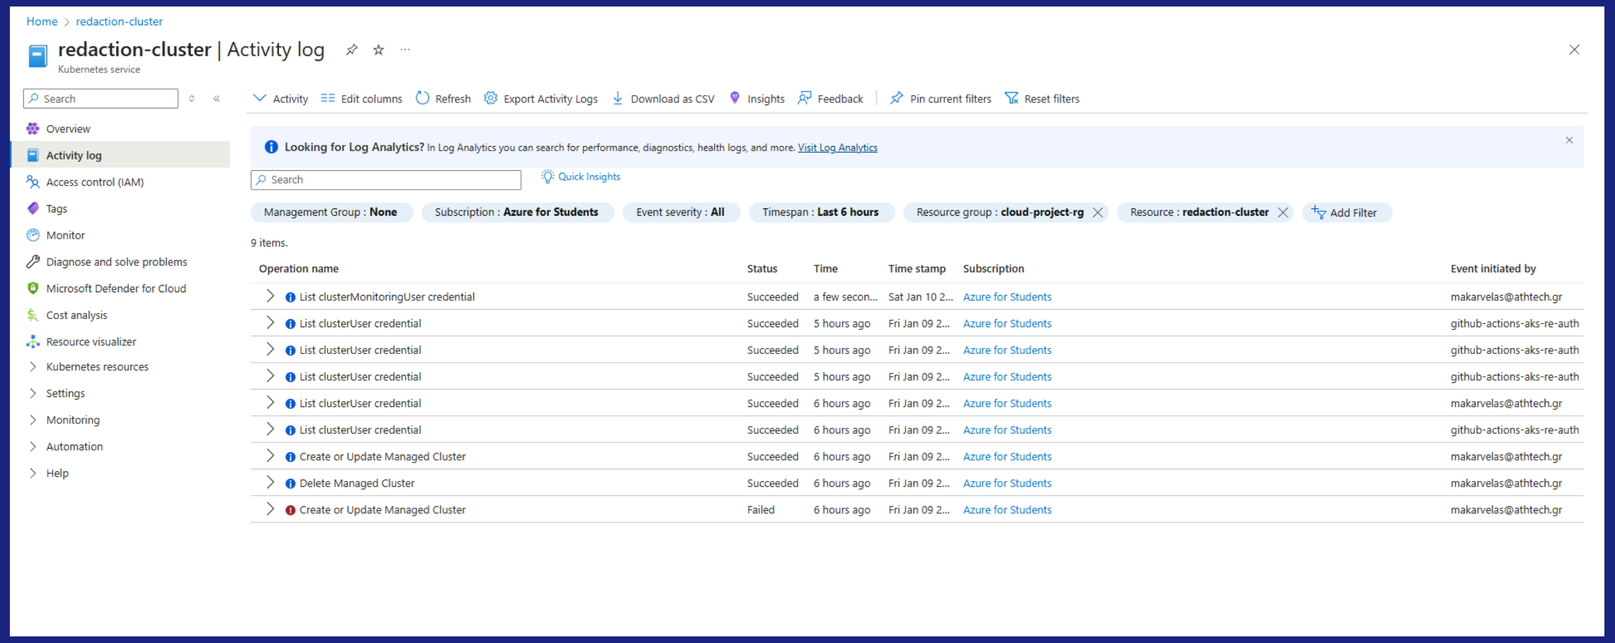

CI/CD Pipeline (DevOps)

- Push: Code committed to GitHub triggers the pipeline.
- Build: Docker Image created and pushed to Azure Container Registry (ACR).
- Deploy: AKS cluster pulls the new image and updates the deployment.

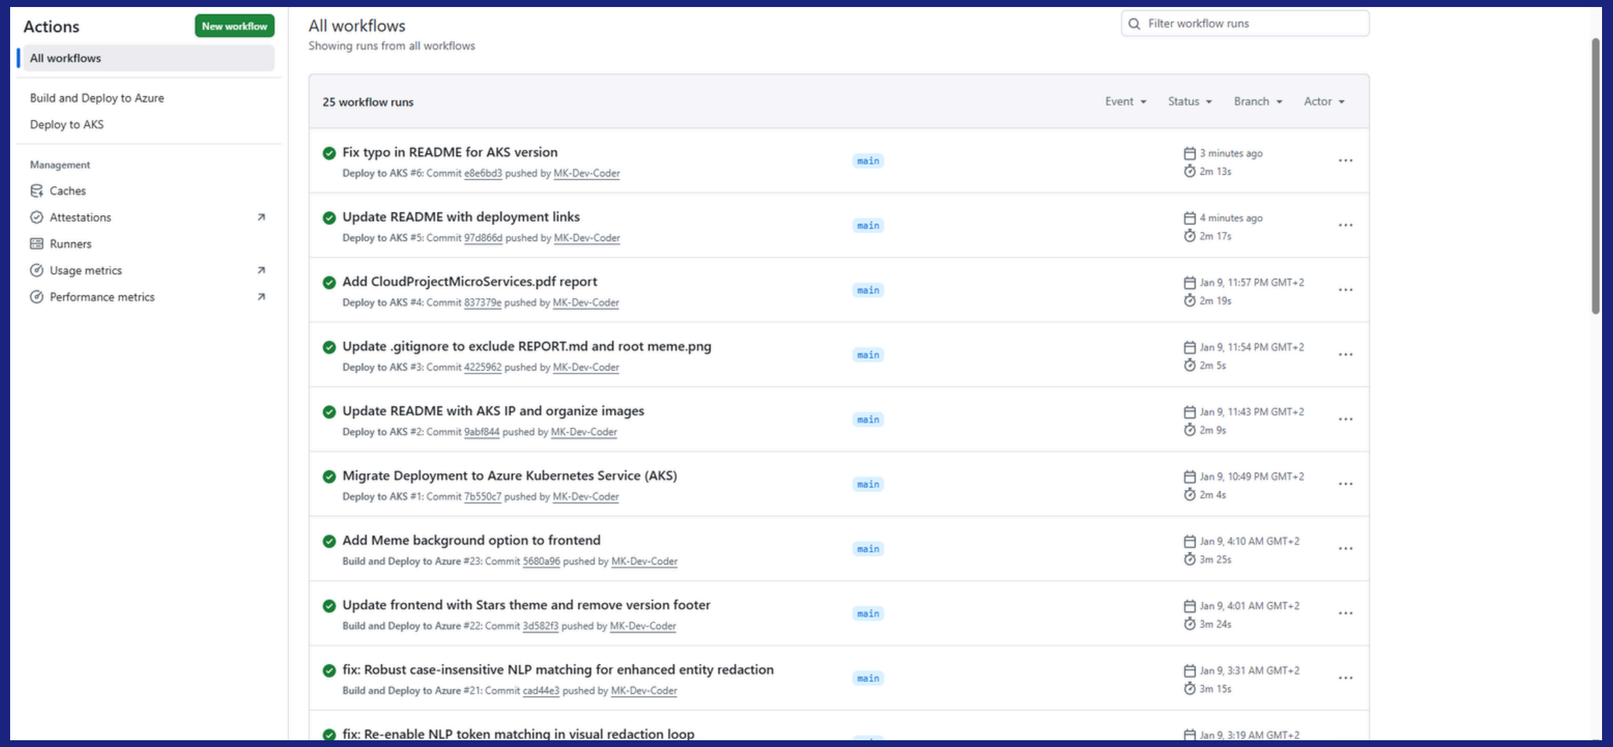

Implementation Steps and Phases

Development began locally, where I built and tested the backend and frontend microservices to ensure stability. After confirming that all components functioned correctly in isolation, I deployed the full solution to Microsoft Azure, making the web application accessible via the public cloud."

Future Improvements

- Advanced Visual Privacy (Scene Text Recognition):
  - Current State: The system works on scanned documents (white background, black text).
  - Future: Implement Object Detection (e.g., using YOLO or CRAFT) to identify and blur sensitive information in natural photos.
  - Use Case: Automatically blurring house numbers, license plates, and street signs in uploaded photos to protect location privacy.
- Multi-Language Support:
  - Future: Integrate additional spaCy models (e.g., es_core_news_sm for Spanish) to support international document redaction.
- User Authentication (RBAC):
- Future: Implement OAuth2 (Login with Google/Microsoft) so that only authorized employees can redact documents, with audit logs tracking who redacted what.

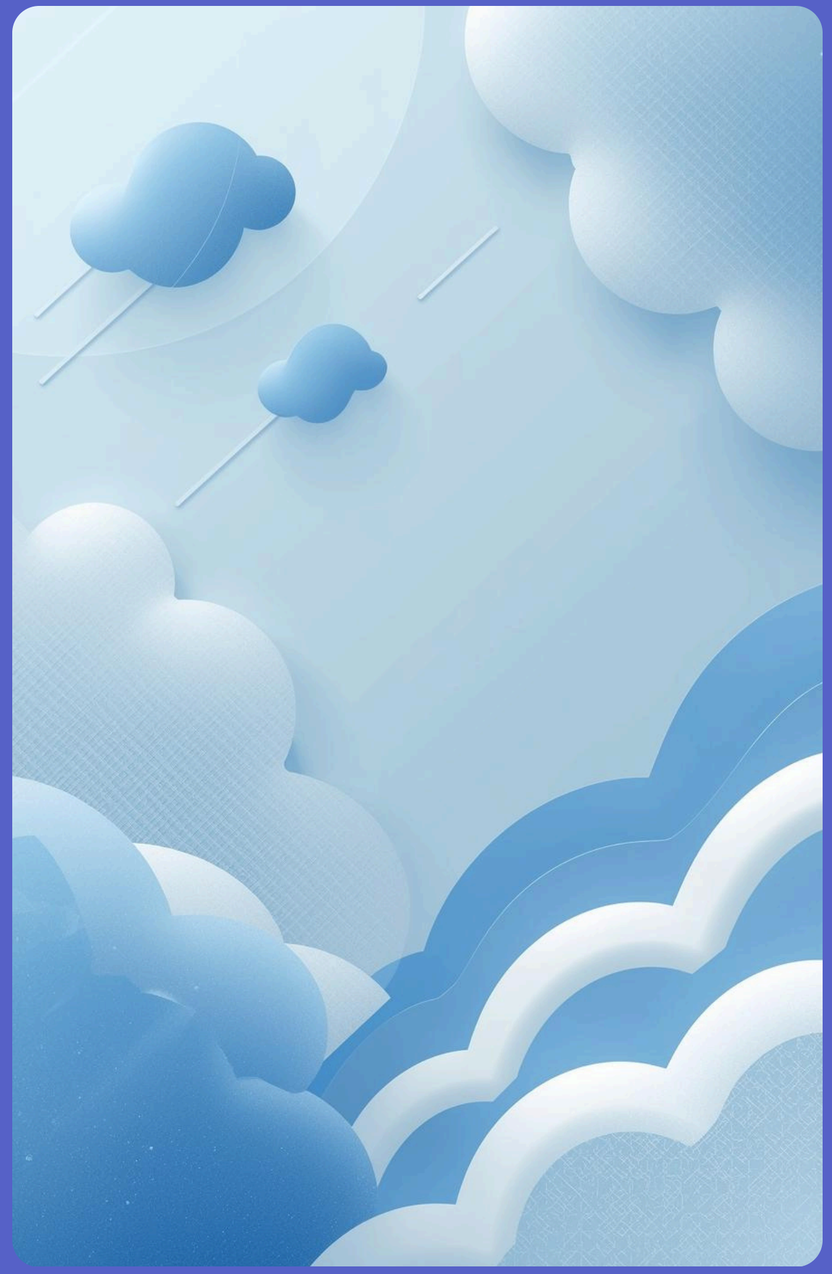

Thank you for your attention